

# SimonsSVM: A Fast and Scalable Support Vector Machine Implementation for R

Philipp Thomann & Ingo Steinwart

Institute of Stochastics and Applications  
University of Stuttgart

European R Users Meeting, Posnan, October 13, 2016

# Support Vector Machines

SVMs are kernel methods for non-parametric classification and regression in the widest sense. They have been studied widely both in theory and in practice.

## Historical Development

**Classification** V. Vapnik et al. (1992 –): Hinge loss

**Regression** G. Wahba (1971 –): Least squares loss

**Other scenarios** since 1995

## Books

Vapnik 1998, Schölkopf/Smola 2002, Steinwart/Christmann 2008.

## Implementations


libsvm (R-package e1071), SVMlight (R-package klaR), kernlab, ...

# SimonsSVM in one Minute

```
install.packages("SimonsSVM",  
                 repos="http://www.isa.uni-stuttgart.de/R")  
library(SimonsSVM)
```

## Least-Squares regression


```
# Load test and training data  
reg <- smldata('reg-1d')  
model <- svm(Y~., reg$train)  
errors(test(model, reg$test))  
#> 0.00541
```



Available regression scenarios: least-squares regression, quantile regression, expectile regression.

## Classification


```
banana <- smldata('banana-bc')  
model <- svm(Y~., banana$train, display=1)  
#> Considering cell 1 out of 1 for task 1 out of 1.  
#> Fold 1: training set size 1600, validation set size 400.  
#> Fold 2: training set size 1600, validation set size 400.  
#> Fold 3: training set size 1600, validation set size 400.  
#> [...]  
errors(test(model, banana$test) )  
#> 0.147
```



Available classification scenarios: multi-class  
classification (default for factors),  
Neyman-Pearson learning, ROC

## Classification

```
banana <- smldata('banana-bc')  
model <- svm(Y~., banana$train, display=1)  
#> Considering cell 1 out of 1 for task 1 out of 1.  
#> Fold 1: training set size 1600, validation set size 400.  
#> Fold 2: training set size 1600, validation set size 400.  
#> Fold 3: training set size 1600, validation set size 400.  
#> [...]  
errors(test(model, banana$test) )  
#> 0.147
```



Available classification scenarios: multi-class classification (default for factors), Neyman-Pearson learning, ROC

# Summary of Features

- Optimized for speed and fully automatic learning.
- Fully integrated cross validation let's the user focus on understandable parameters.
- Meaningful default values often make tuning of parameters unnecessary.
- Partitioning strategies decrease memory usage and training and testing time for big problems. Large-scale problems become treatable.
- Many standard learning scenarios are included.
- Comprehensive documentation.
- Kernel computation is multi-threaded and can make use of GPUs.
- Provides Gauss RBF and Poisson kernel, more can be added.
- Bindings for R/MATLAB/Java/Python
- Open Source

# Benchmarks

## Test Errors (relative to SimonsSVM)

	Single runs with the recommended param.				5-fold Cross-valid. on $10 \times 10$ -grid	
	SimonsSVM-cv	libsvm	SVMlight	kernlab	SimonsSVM-cv	libsvm-cv
BANK-MARKETING	1.0	1.089	1.089	1.085	1.0	0.990
COD-RNA	1.0	2.878	1.428	1.342	1.0	0.931
COVTYPE	1.0	1.351	1.353	1.213	1.0	0.999
THYROID-ANN	1.0	1.529	1.533	1.377	1.0	0.949

Always mean of 10 single-threaded classification-runs on samples of size 2000 on an Intel® Core®-i7 with 16GB RAM running Debian Linux.

## Runtime (relative to SimonsSVM)

	Single runs with the recommended param.			
	SimonsSVM	libsvm	SVMlight	kernlab
BANK-MARKETING	1.0	2.8	6.1	3.4
COD-RNA	1.0	6.0	5.0	3.2
COVTYPE	1.0	8.1	12.2	5.9
THYROID-ANN	1.0	2.2	5.7	3.6

# Benchmarks

## Test Errors (relative to SimonsSVM)

	Single runs with the recommended param.				5-fold Cross-valid. on $10 \times 10$ -grid	
	SimonsSVM-cv	libsvm	SVMlight	kernlab	SimonsSVM-cv	libsvm-cv
BANK-MARKETING	1.0	1.089	1.089	1.085	1.0	0.990
COD-RNA	1.0	2.878	1.428	1.342	1.0	0.931
COVTYPE	1.0	1.351	1.353	1.213	1.0	0.999
THYROID-ANN	1.0	1.529	1.533	1.377	1.0	0.949

Always mean of 10 single-threaded classification-runs on samples of size 2000 on an Intel® Core®-i7 with 16GB RAM running Debian Linux.

## Runtime (relative to SimonsSVM)

	Single runs with the recommended param.			
	SimonsSVM	libsvm	SVMlight	kernlab
BANK-MARKETING	1.0	2.8	6.1	3.4
COD-RNA	1.0	6.0	5.0	3.2
COVTYPE	1.0	8.1	12.2	5.9
THYROID-ANN	1.0	2.2	5.7	3.6



# Benchmarks

## Test Errors (relative to SimonsSVM)

	Single runs with the recommended param.				5-fold Cross-valid. on $10 \times 10$ -grid	
	SimonsSVM-cv	libsvm	SVMlight	kernlab	SimonsSVM-cv	libsvm-cv
BANK-MARKETING	1.0	1.089	1.089	1.085	1.0	0.990
COD-RNA	1.0	2.878	1.428	1.342	1.0	0.931
COVTYPE	1.0	1.351	1.353	1.213	1.0	0.999
THYROID-ANN	1.0	1.529	1.533	1.377	1.0	0.949

Always mean of 10 single-threaded classification-runs on samples of size 2000 on an Intel® Core®-i7 with 16GB RAM running Debian Linux.

## Runtime (relative to SimonsSVM)


	Single runs with the recommended param.				5-fold Cross-valid. on $10 \times 10$ -grid	
	SimonsSVM	libsvm	SVMlight	kernlab	SimonsSVM-cv	libsvm-cv
BANK-MARKETING	1.0	2.8	6.1	3.4	27.5	2310.4
COD-RNA	1.0	6.0	5.0	3.2	24.0	1789.3
COVTYPE	1.0	8.1	12.2	5.9	29.8	3197.9
THYROID-ANN	1.0	2.2	5.7	3.6	27.8	1623.9

# Split Feature space into Cells

SVMs usually are at least  $O(n^2)$  in both time and memory. Simons' SVM has a solution

```
svm(Y~., covtype_full, useCells=TRUE)
```

You can choose from 6 different methods.




The full COVERTYPE data set (about 523.000 samples) would need 1.6TB of memory. Simons' SVM, by splitting it into cells of size at most 2000, trains and tests on a Core-i5 laptop with 12GB of memory in 9.8 minutes!

# Split Feature space into Cells

SVMs usually are at least  $O(n^2)$  in both time and memory. Simons' SVM has a solution

```
svm(Y~., covtype_full, useCells=TRUE)
```


You can choose from 6 different methods.



The full COVERTYPE data set (about 523.000 samples) would need 1.6TB of memory. Simons' SVM, by splitting it into cells of size at most 2000, trains and tests on a Core-i5 laptop with 12GB of memory in 9.8 minutes!

# Split Feature space into Cells (II)

This works without sacrificing generalization:



Also large-scale data sets are possible to treat:

E.g. we trained and tested the HIGGS data set (10 million samples) in five hours on a Xeon hexacore with 64GB memory.

# Implementation Details

- Written in C++ and its main functionality is accessible through a small number of high-level C++ classes.
- The routines for computing the kernel matrices and for evaluating the SVM models on the test data are multi-threaded.
- Time critical inner loops are vectorized whenever this is possible. (SSE2, AVX, and AVX2)

The solution algorithms are of SMO-type. They do not use any offset in the decision functions: This is not needed for Gaussian RBF and Poisson kernels, and it makes the algorithm much faster and better to use in cross-validation.

# Interfaces

There is a command line version of Simons' SVM with scripts for the learning scenarios.

The R-bindings are the most mature of the bindings:

- Full integration into the R-workflow including several helper functions
- 2 Vignettes: demo and documentation
- online help (e.g. `?svm`, ...)

All the bindings (R/MATLAB/Java/Python):


- share a common C-interface to Simons' SVM C++ code.
- can be compiled either optimized for the current machine (`native`) or to use only SSE2 (`generic`).
- All operations are in-memory and in-process.

The command line version and all the bindings have been tested on Linux and macOS, as well as on Windows 8.

# Multiclass classification

Multiclass classification has to be reduced to binary tasks:

- all-vs-all (Default)
- one-vs-all (aka one-vs-rest)




```
model <- svm(Y~., banana_mc$train)
# banana_mc$train$Y is factor with 4 levels
errors(test(model, banana_mc$test))
#> result 1vs2 1vs3 1vs4 2vs3 2vs4 3vs4
#> 0.217 0.142 0.112 0.093 0.074 0.074 0.001
model <- mcSVM(Y~., banana_mc$train, mc_type="OvA_ls")
errors(test(model, banana_mc$test))
#> result 1vsOthers 2vsOthers 3vsOthers 4vsOthers
#> 0.2147 0.1545 0.1227 0.0777 0.0737
```

The first element in the errors is the overall test error.

## Quantile regression

This uses the quantile solver with pinball loss and performs selection for every quantile provided.

```
quantiles_list <- c(0.05, 0.1, 0.5, 0.9, 0.95)  
model <- qtSVM(Y ~ ., reg$train, weights=quantiles_list)
```




In this plot you see estimations for two lower and upper quantiles as well as the median of the distribution of the label  $y$  given  $x$ .



## Expectile regression (solver implemented by M. Farooq)

This uses the expectile solver with weighted least squares loss and performs selection for every weight. The 0.5-expectile in fact is just the ordinary least squares regression and hence estimates the mean of  $y$  given  $x$ . And in the same way as quantiles generalize the median, expectiles generalize the mean.

```
expectiles_list <- c(0.05, 0.1, 0.5, 0.9, 0.95)  
model <- exSVM(Y ~ ., reg$strain, weights=expectiles_list)
```



# Neyman-Pearson-Learning

Neyman-Pearson-Learning attempts classification under the constraint that the probability of false positives (Type-I error) is bound by a significance level  $\alpha$ , which we call the NPL-constraint.


```
model <- nplSVM(Y ~ ., banana, class=-1)
# [...]
#>
#>           1      2      3      4      5
#> npl_constraints 0.025 0.0333 0.050 0.075 0.100
#> false_alarm_rate 0.037 0.0390 0.053 0.058 0.086
#> detection_rate  0.563 0.5630 0.678 0.692 0.770
```

You can see that the false alarm rate in the test set meet the NPL-constraints quite nicely, and on the other hand the the detection rate is increasing.

# ROC curve

Receiver Operating Characteristic curve (ROC curve) plots trade-off between the false alarm rate and the detection rate for different weights.

```
model <- rocSVM(Y ~ ., banana)
```



This shows nice learning, since the ROC curve is near the north-west corner.

# Questions? Bug-reports?

Please visit our Homepage

`http://www.isa.uni-stuttgart.de/software/`

or write me an E-Mail

`philipp.thomann@mathematik.uni-stuttgart.de`

# Questions? Bug-reports?

Please visit our Homepage

`http://www.isa.uni-stuttgart.de/software/`

or write me an E-Mail

`philipp.thomann@mathematik.uni-stuttgart.de`

Dziękuję