

brms: An R Package for Bayesian Multilevel Models using Stan

Paul Bürkner

Institut für Psychologie
Westfälische Wilhelms-Universität Münster

13.10.2016



Agenda

1 Short introduction to Stan

2 The brms package

- Model Specification
- Model Fitting
- Post-Processing

3 Discussion



Stan Overview

- Probabilistic programming language...
- ... to fit open-ended Bayesian Models
- Powerful sampling algorithms: HMC and NUTS
- Automatic differentiation library
- Runs on all major platforms (Windows, OS X, Linux)
- Can be called from R, Python, Julia, Stata, and Matlab



Stan code for simple linear regression

```
1 data {
2   int<lower=1> N; // number of observations
3   vector[N] Y; // response variable
4   int<lower=1> K; // number of fixed effects
5   matrix[N, K] X; // FE design matrix
6 }
7 parameters {
8   vector[K] b; // fixed effects
9   real<lower=0> sigma; // residual SD
10 }
11 transformed parameters {
12   vector[N] eta; // linear predictor
13   eta <- X * b;
14 }
15 model {
16   // prior specifications
17   sigma ~ student_t(3, 0, 100);
18   // likelihood contribution
19   Y ~ normal(eta, sigma);
20 }
```



The idea of brms

Make Stan usable for a wider range of audience

- use R formula syntax
- internally write Stan code that is readable yet fast
- provide an easy interface for defining priors
- facilitate post-processing



Some features of brms (1)

Basic model types:

- (Robust multivariate) linear models
- Count data models
- Categorical and ordinal models
- Survival models
- Zero-inflated and hurdle models
- Non-linear models
- ...



Some features of brms (2)

Other modeling options:

- Group-level terms using **lme4** syntax
- smoothing terms
- ordinal predictors
- autocorrelation
- meta-analytic standard errors
- censored / truncated data
- ...



Example: Censored survival regression

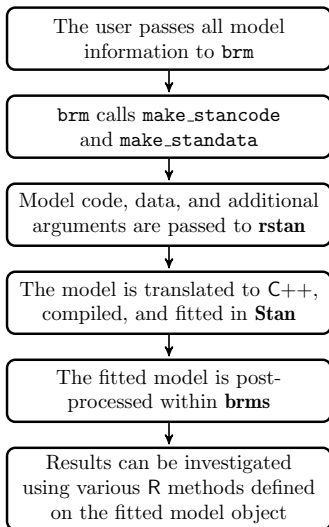
```
1 | brm(time | cens(censored) ~ age * sex + (1 + age|patient),  
2 |     data = kidney, family = exponential())
```

Let's add some priors:

```
1 | brm(time | cens(censored) ~ age * sex + (1 + age|patient),  
2 |     data = kidney, family = exponential(),  
3 |     prior = c(set_prior("normal(0,5)", class = "b"),  
4 |               set_prior("cauchy(0,2)", class = "sd"),  
5 |               set_prior("lkj(2)", class = "cor")))
```




Flow chart of the model fitting procedure





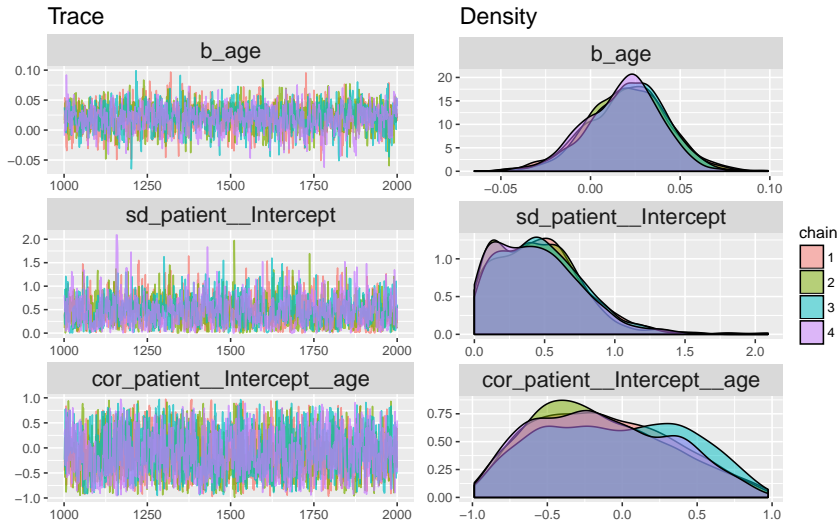
Example: Summary

```

1 Family: exponential (log)
2 Formula: time | cens(censored) ~ age * sex + (1 + age | patient)
3 Data: kidney (Number of observations: 76)
4 Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
5           total post-warmup samples = 4000
6 WAIC: 668.22
7
8 Random Effects:
9 ~patient (Number of levels: 38)
10
11           Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
12 sd(Intercept)      0.45      0.30      0.02      1.07      1172 1.00
13 sd(age)            0.01      0.01      0.00      0.03       450 1.01
14 cor(Intercept,age) -0.12      0.45     -0.87      0.77      1133 1.00
15
16 Fixed Effects:
17           Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
18 Intercept          2.91      0.93      1.22      4.85      1636 1
19 age                0.02      0.02     -0.02      0.06      1527 1
20 sexfemale          2.76      1.07      0.59      4.77      1578 1
21 age:sexfemale      -0.03      0.02     -0.08      0.02      1478 1


```

Example: Trace and density plots





Example: Marginal plots






Example: Hypothesis Testing

```
1 > hypothesis(fit, "Intercept - age > 0", class = "sd", group = "patient")
2 Hypothesis Tests for class sd_patient:
3       Estimate Est.Error l-95% CI u-95% CI Evid.Ratio
4 Intercept-age > 0      0.44      0.3      0.03      Inf      62.49 *
```

Hypothesis for class sd_patient





Plans for the future

- Full support for multivariate models
- Spatial correlation structures
- Mixture distributions
- Error in variables models
- Missing data imputation
- ...

Your ideas are always welcome!



Learn more about brms and Stan

- Help within R: `vignette("brms")` or `help("brms")`
- List of all applicable methods: `methods(class = "brmsfit")`
- Website of brms: <https://github.com/paul-buerkner/brms>
- Website of Stan: <http://mc-stan.org/>
- Contact me: paul.buerkner@gmail.com



Thank you for your attention!

Questions?